

# Basic SQL

# Introduction

- We deal with a **large amount** of data. Need a way to **define** the database, tables, etc. **Process** data more effectively. **Turn** the data into information.
- And this is where **SQL** comes to play!
  - The **most commonly used language** to interface with a DB system is the **Structured Query Language (SQL)**
- MySQL is a **database management system**
  - Allows you to **manage** relational databases
  - Open source! by Oracle

# What is MySQL?

- We deal with a **large amount** of data. Need a way to **define** the database, tables, etc. **Process** data more effectively. **Turn** the data into information.
- And this is where **SQL** comes to play!
- MySQL is a **database management system**
  - Allows you to **manage** relational databases
  - Open source! by Oracle



# Install MySQL

- Download MySQL installer from:  
<http://dev.mysql.com/downloads/installer/>
- A tutorial shows you step by step how to install MySQL using MySQL Installer: <http://www.mysqltutorial.org/install-mysql/>



# MySQL Workbench

- A unified visual tool for database architects, developers, and DBAs
- <https://www.mysql.com/products/workbench/>



# Introduction

- You have learned how to formulate queries in a mathematical language (relational algebra), it is time to put it into **practice**
- SQL is based on relational algebra

# Introduction (cont.)

- A "case-in-point" Example: Find FName, LName of **employees** in the "Research" **department**

- Relational algebra:

$$\pi_{fname, lname}(Employee \bowtie_{dno = dnumber} \sigma_{dname = 'Research'}(Department))$$

$$\pi_{fname, lname}(\sigma_{dno = dnumber \text{ AND } dname = 'Research'}(Employee \times Department))$$

- SQL:

```
SELECT fname, lname
FROM   employee, department
WHERE  dno=dnumber
      AND dname='Research'
```

- I don't know about you, but I can clearly see the correlation between Relational Algebra and SQL:
  1. The **FROM** clause specifies a **Cross product**
  2. The **WHERE** clause specifies a selection operation
  3. The **SELECT** clause specifies a projection operation

# Introduction (cont.)

- Relational algebra are **low-level** language for DBMS users
  - A sequence of operations
  - User must specify **how**—that is, *in what order* to execute
- This chapter presents practical relational model
  - Based on SQL standard
  - The **SQL** language provides a **high-level** *declarative* language interface
  - User only specifies **what** the result is to be
  - More user-friendly



# Introduction (cont.)

- SQL is **not only** a **query language** - i.e., a language used to **retrieve information** from a database
- SQL also contain:
  - a **Data Definition Language** - i.e.:
    - a language that is used to **define the database and its objects**, e.g., tables, views, etc.
  - a **Data Manipulation Language** - i.e.:
    - a language that is used to **update and query data**
  - a **Data Control Language**
    - a language that is used to **grant the permissions** to a user to access a certain data in the database

# Order of Discussion of Chp 6 and Chp 7

- Data Definition
- Retrieval Query (SELECT)
- Data update (INSERT, DELETE, UPDATE)
- View definition

# Today's Lecture

1. SQL for creating schemas and tables
2. An overview of basic data types in SQL
3. How basic constraints are specified
4. Specify retrieval queries (SELECT)
5. SQL commands for insertion, deletion, and update

# Data Definition Language Features in SQL

- SQL allow users to perform the following **data definition functions**:
  - **Create** a database
  - **Define** new relations in a database
  - **Define** constraints on attributes in the relations
  - **Alter** the **structure** of (existing) relations  
(This operation can be very complex when there is already data in the relations - it is best done when relations are still empty)
  - **Delete (drop)** relations

# Creating a Database: the CREATE SCHEMA command

- The **CREATE SCHEMA** command is used to *define a database schema*
  - The *description* of a database is called the **database schema**
    - Used to group together database tables (= relations)
    - Also contains other constructs (such as constraints)

**Figure 2.1**

Schema diagram for the database in Figure 1.2.

## STUDENT

|      |                |       |       |
|------|----------------|-------|-------|
| Name | Student_number | Class | Major |
|------|----------------|-------|-------|

## COURSE

|             |               |              |            |
|-------------|---------------|--------------|------------|
| Course_name | Course_number | Credit_hours | Department |
|-------------|---------------|--------------|------------|

## PREREQUISITE

|               |                     |
|---------------|---------------------|
| Course_number | Prerequisite_number |
|---------------|---------------------|

## SECTION

|                    |               |          |      |            |
|--------------------|---------------|----------|------|------------|
| Section_identifier | Course_number | Semester | Year | Instructor |
|--------------------|---------------|----------|------|------------|

## GRADE\_REPORT


|                |                    |       |
|----------------|--------------------|-------|
| Student_number | Section_identifier | Grade |
|----------------|--------------------|-------|

# Creating a Database: the CREATE SCHEMA command (cont.)

- Syntax of the CREATE SCHEMA command in the SQL standard:

**CREATE SCHEMA** schema\_name **AUTHORIZATION** db\_user;

**CREATE SCHEMA** COMPANY **AUTHORIZATION** 'Jsmith';



To indicate  
who own the  
schema

- When a database schema is created, **ALL access authorization** is granted to the user **db\_user** (i.e., he is the **owner** of the database)

# Creating a Database: the CREATE SCHEMA command (in MySQL)

- **MySQL** version of **Create Schema**

```
create database database_name ;
```

- The **database** is created by the **root user**
- The **authorization** is granted **separately**
- **See example:**

# The CREATE TABLE Command in SQL

- The **CREATE TABLE** command is used to *define* tables ( = relations) within a database schema
- The **relation** created by the **CREAT TABLE** command is initially *empty*
- You can then use an **INSERT** command to *insert data into the relation*



Th CREATE TABLE EMPLOYEE ( Fname VARCHAR(15) NOT NULL, t.)

• Sy

CRI

(

~

~

CREATE TABLE DEPARTMENT

~

[

);

( Fname VARCHAR(15) NOT NULL,  
Minit CHAR,  
Lname VARCHAR(15) NOT NULL,  
Ssn CHAR(9) NOT NULL,  
Bdate DATE,  
Address VARCHAR(30),  
Sex CHAR,  
Salary DECIMAL(10,2),  
Super\_ssn CHAR(9),  
Dno INT NOT NULL,

PRIMARY KEY (Ssn),

( Dname VARCHAR(15) NOT NULL,  
Dnumber INT NOT NULL,  
Mgr\_ssn CHAR(9) NOT NULL,  
Mgr\_start\_date DATE,

PRIMARY KEY (Dnumber),  
UNIQUE (Dname),  
FOREIGN KEY (Mgr\_ssn) REFERENCES EMPLOYEE(Ssn) );

# Data Type in SQL

- Each attribute has a **data type**
- Numeric types:
  - **INTEGER** or **INT**
  - **DECIMAL(*i*, *j*)**
    - DEC(8,3) = fixed point numbers of the format: xxxxx.yyy
- Character string
  - **CHARACTER(*n*)** or **CHAR(*n*)**
    - Fixed length character strings (number of characters in string is (always) *n*)
  - **VARCHAR(*n*)**
    - Variable length character strings (maximum number of characters in string is *n*)
  - Ex.: To denote string, use **single quotes**: 'John'
- Bit string
  - **BIT(*n*)**
    - Fixed length bit string (number of bits in string is (always) *n*)
  - **BIT VARYING(*n*)**
    - Variable length bit string (maximum number of bits in string is *n*)
  - To denote bit string 10101, use: **B'10101'**

# Data Type in SQL (cont.)

- Boolean
  - **TRUE, FALSE** or **UNKNOWN**
    - if either x or y is NULL , then some logical comparisons evaluate to an UNKNOWN
- Date
  - DATE values are specified as: **DATE'YYYY-MM-DD'**
  - Must be preceded by the keyword DATE
- Time
  - TIME values are specified as: **TIME'HH:MM:SS'**
  - Must be preceded by the keyword TIME

# Specifying Constraints on Relations

- The **CREATE TABLE** command can be used to *specify*:
  - Constraints
  - Default values
    - Specify default values for attributes
- Types of constraints
  - Primary Key constraint
  - Entity Integrity
  - Referential integrity constraint
  - Domain constraint
  - Not-NULL constraint

# Specifying Primary Key Constraints

- To **define** that a given set of **attribute** constitute the **primary key** of the relation, we use:
  - **PRIMARY KEY**(attribute-list), or
  - Dnumber INT **PRIMARY KEY** (if a primary key has a **single** attribute)

```
CREATE TABLE EMPLOYEE
```

|                           |                |           |
|---------------------------|----------------|-----------|
| ( Fname                   | VARCHAR(15)    | NOT NULL, |
| Minit                     | CHAR,          |           |
| Lname                     | VARCHAR(15)    | NOT NULL, |
| Ssn                       | CHAR(9)        | NOT NULL, |
| Bdate                     | DATE,          |           |
| Address                   | VARCHAR(30),   |           |
| Sex                       | CHAR,          |           |
| Salary                    | DECIMAL(10,2), |           |
| Super_ssn                 | CHAR(9),       |           |
| Dno                       | INT            | NOT NULL, |
| <b>PRIMARY KEY (Ssn),</b> |                |           |

# Specifying Primary Key Constraints (cont.)

- Example: **multiple** attributes in key

```
CREATE TABLE WORKS_ON
( Essn          CHAR(9)          NOT NULL,
  Pno           INT              NOT NULL,
  Hours        DECIMAL(3,1)     NOT NULL,
  PRIMARY KEY (Essn, Pno),
```

- The third will fail:

```
insert into test2 values ('111223333', 44, 4.5);
insert into test2 values ('111223333', 23, 3.5);
insert into test2 values ('111223333', 44, 6.5);
```

# Specifying Primary Key Constraints (cont.)

- There may be **more than one keys** in a relation
- The UNIQUE constraints can be used to specify **candidate keys**:
  - **UNIQUE**(attribute-list)

```
CREATE TABLE DEPARTMENT
( Dname          VARCHAR(15)          NOT NULL,
  Dnumber        INT                  NOT NULL,
  Mgr_ssn        CHAR(9)              NOT NULL,
  Mgr_start_date DATE,
  PRIMARY KEY (Dnumber),
  UNIQUE (Dname),
```

- Now the second insert will **fail**:
  - insert into test3 values (4, 'Research')
  - insert into test3 values (5, 'Research')

# Specifying Referential Integrity Constraints

- Recall that:

A **Foreign Key** is a set of attributes used to reference (identify) **tuples in another relation**

The **referential integrity constraint** states that the **referenced tuples must exist**

- To define that a given set of attribute constitute a foreign key, we use:

- **FOREIGN KEY** (attribute-list) **REFERENCES** relation(attribute-list)

```
CREATE TABLE DEPARTMENT
```

|                |             |           |
|----------------|-------------|-----------|
| ( Dname        | VARCHAR(15) | NOT NULL, |
| Dnumber        | INT         | NOT NULL, |
| Mgr_ssn        | CHAR(9)     | NOT NULL, |
| Mgr_start_date | DATE,       |           |

```
PRIMARY KEY (Dnumber),
```

```
UNIQUE (Dname),
```

```
FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn) );
```



# Another Example

```
CREATE TABLE test4
(
    ssn CHAR(9),                /* <----- key */
    salary dec(9,2),

    CONSTRAINT Test4PrimKey PRIMARY KEY(ssn)
);

CREATE TABLE test5
(
    essn CHAR(9),
    pno INTEGER,

    CONSTRAINT Test5ForeignKey
        FOREIGN KEY (essn)
        REFERENCES test4(ssn) /* Used as foreign key */
);
```

- Now this insert will **fail**:
  - Insert into test5 values ('111223333', 5);
  - Because there is **no tuple** in test4 with SSN='111223333'
- To make this insert legal, we must **first** insert the employee:
  - Insert into test4 values ('111223333', 5000.00);

# Specifying Referential Integrity Constraints (cont.)

- The attributes used as a foreign key must be specified as:
  - Primary key, or as
  - Unique
- Example: this will cause an **error**:

```
CREATE TABLE test4
(
    ssn CHAR(9),           // Not primary key nor UNIQUE
    salary dec(9,2)
);

CREATE TABLE test5
(
    essn CHAR(9),
    pno INTEGER,

    CONSTRAINT Test5ForeignKey
    FOREIGN KEY (essn)
    REFERENCES test4(ssn)
);
```

# Given Names to Constraints

- A **constraint name** is used to **identify** a particular constraint in case the constraint **must be dropped later** and **replaced** with another constraint,
- Example:
  - **CONSTRAINT** [constraint-name] **PRIMARY KEY**(attribute-list)
  - **CONSTRAINT** [constraint-name] **UNIQUE**(attribute-list)
  - **CONSTRAINT** [constraint-name]  
    **FOREIGN KEY** (attribute-list)  
    **REFERENCES** relation(attribute-list)

Giving names to constraints is **optional**

# Specifying Constraints on Attributes

- SQL also allow the user to define a number of constraints on attribute values (domain constraints)
  - Not Null
    - Specifies that the attribute value **cannot** have the **NULL** value
  - Default value
  - Domain range
- Example:

```
CREATE TABLE test6
(  
    ssn CHAR(9) NOT NULL,  
    fname CHAR(30),  
    lname CHAR(30)  
);
```

- This insert will **succeed**:
  - insert into test6 values ('123456789', 'John', 'Smith');
- But this insert will **fail**:
  - insert into test6 values (null, 'John', 'Smith');

# Specifying Constraints on Attributes (cont.)

- Specifies a **default value** for an attribute

```
CREATE TABLE test7  
(  
    ssn CHAR(9) NOT NULL,  
    salary DECIMAL(6,2) DEFAULT 5000  
);
```

- This is a insert using a **partial set of attributes**:
  - insert into test7(ssn) values ('111223333');
  - The tuple does not contain the **salary attribute**
- The default value 5000 is that assigned to the salary attribute , so effectively:  
insert into test7 values ('111223333', 5000);

# Specifying Constraints on Tuples Using CHECK

- Row-based Constraints
  - Apply to each row individually and are checked whenever a row is inserted or modified
- Example

```
CREATE TABLE test8
(  
    ssn CHAR(9) NOT NULL,  
    dno INTEGER CHECK (dno > 0 and dno < 21)  
);
```

- This insert will **succeed**:
  - insert into test8 values ('111223333', 11);
- This insert will **fail**:
  - insert into test8 values ('333333333', 44);

# SQL: SELECT

- The SQL SELECT Command
  - Is used to **retrieve** the set of **tuples** that **satisfy** a given **condition**
  - Is based on **Relational Algebra** and you will see the **join/section/projection** operations reflected in the SELECT command
- The basic form of the SELECT command
  - SELECT-FROM-WHERE block:

```
SELECT attribute_list
FROM   relation_list
WHERE  boolean_expression
```

1. **attribute\_list**: a list of attributes **from** the relations in the relation\_list
2. **relation\_list**: a list of relations. The **CROSS PRODUCT** of these relations is formed
3. **boolean\_expression**: condition

# SQL: SELECT (cont.)

```
SELECT attr1, attr2, ..., attrM  
FROM    R1, R2, R3, ..., RN  
WHERE   boolean_expression
```

- Meaning of the SQL SELECT command:
  - First, **form** the **Cross product** of the relations  $R1 \times R2 \times R3 \times \dots \times RN$
  - Next, **select** all the **tuples** that **satisfy** the **boolean\_expression**
  - Finally, **project** the **attributes** attr1, attr2, ..., attrM from the **qualifying tuples**



# Running Examples

- Example 1: **List** SSN, Lname and DNO of all **employees**

```
select ssn, lname, dno  
from employee
```

| SSN       | LNAME   | DNO |
|-----------|---------|-----|
| 123456789 | Smith   | 5   |
| 333445555 | Wong    | 5   |
| 999887777 | Zelaya  | 4   |
| 987654321 | Wallace | 4   |
| 666884444 | Narayan | 5   |
| 453453453 | English | 5   |
| 987987987 | Jabbar  | 4   |
| 888665555 | Deena   | 1   |

```
select dnumber, dname  
from department
```

| DNUMBER | DNAME |
|---------|-------|
|---------|-------|

|   |                |
|---|----------------|
| 5 | Research       |
| 4 | Administration |
| 1 | Headquarters   |

- Exam

# Running Examples

- Example 3: See **what happens** when “Department” are specified in the r

Observation:

- The **output** is the **cartesian product** of the two relations "employee" and "department"

Query:

```
select ssn, lname, dno, dnumber, dname
from employee, department
```

Output:

| SSN       | LNAME   | DNO | DNUMBER | DNAME          |
|-----------|---------|-----|---------|----------------|
| 123456789 | Smith   | 5   | 5       | Research       |
| 333445555 | Wong    | 5   | 5       | Research       |
| 999887777 | Zelaya  | 4   | 5       | Research       |
| 987654321 | Wallace | 4   | 5       | Research       |
| 666884444 | Narayan | 5   | 5       | Research       |
| 453453453 | English | 5   | 5       | Research       |
| 987987987 | Jabbar  | 4   | 5       | Research       |
| 888665555 | Borg    | 1   | 5       | Research       |
| 123456789 | Smith   | 5   | 4       | Administration |
| 333445555 | Wong    | 5   | 4       | Administration |
| 999887777 | Zelaya  | 4   | 4       | Administration |
| 987654321 | Wallace | 4   | 4       | Administration |
| 666884444 | Narayan | 5   | 4       | Administration |
| 453453453 | English | 5   | 4       | Administration |
| 987987987 | Jabbar  | 4   | 4       | Administration |
| 888665555 | Borg    | 1   | 4       | Administration |
| 123456789 | Smith   | 5   | 1       | Headquarters   |
| 333445555 | Wong    | 5   | 1       | Headquarters   |
| 999887777 | Zelaya  | 4   | 1       | Headquarters   |
| 987654321 | Wallace | 4   | 1       | Headquarters   |
| 666884444 | Narayan | 5   | 1       | Headquarters   |
| 453453453 | English | 5   | 1       | Headquarters   |
| 987987987 | Jabbar  | 4   | 1       | Headquarters   |
| 888665555 | Borg    | 1   | 1       | Headquarters   |

# Running Examples

- Example 4: **demonstrates** how to perform a **join** operation on "employee" and "department"

Query:

```
select ssn, lname, dno, dnumber, dname
from employee, department
where dno = dnumber
```

Output:

| SSN       | LNAME   | DNO | DNUMBER | DNAME          |
|-----------|---------|-----|---------|----------------|
| 888665555 | Borg    | 1   | 1       | Headquarters   |
| 999887777 | Zelaya  | 4   | 4       | Administration |
| 987654321 | Wallace | 4   | 4       | Administration |
| 987987987 | Jabbar  | 4   | 4       | Administration |
| 123456789 | Smith   | 5   | 5       | Research       |
| 453453453 | English | 5   | 5       | Research       |
| 666884444 | Narayan | 5   | 5       | Research       |
| 333445555 | Wong    | 5   | 5       | Research       |

# Running Examples

- Example 4: **demonstrates** how to perform a **join** operation on "**employee**" and "**department**"
  - For each tuple (row):

**dno == dnumber** (as according to the where condition)

- In fact, this is what we have learned before in **Relational Algebra**:

$$R_1 \bowtie_{condition} R_1 = \sigma_{condition}(Department)$$

- A **join** operation is a **Cross product** followed by a **selection** operation
- The **FROM** clause in the SQL command specifies the **Cross product** operation
- The **WHERE** clause in the SQL command specifies the **condition of the  $\sigma$**  operation