

SQL: SELECT (cont.)

- **Example 1:** find the name and address of employees working in the 'Research' department

```
SELECT fname, lname, address  
FROM Employee, Departemnt  
WHERE dno = dnumber AND dname = 'Research'
```

- The following is the equivalence query in Relational Algebra:

▪ $\pi_{\text{fname, lname, address}} (\sigma_{(\text{dname}='Research' \text{ and } \text{dnumber} = \text{dno})} (\text{employee} \times \text{department}))$

SQL: SELECT (cont.)

- **Example 2:** For the projects located in 'Stafford', find the name of the project, the name of the controlling department, the last name of the department's manager, his address and birthdate

Projects in 'Stafford' can be found as follows:

```
SELECT pnumber, pname, plocation  
FROM project  
WHERE plocation = 'Stafford';
```

The controlling department of these projects is added as follows:

```
SELECT pnumber, pname, dname, mgrssn  
FROM project, department  
WHERE dnum = pnumber AND plocation = 'Stafford';
```

Finally, the manager information is added as follows:

```
SELECT pname, dname, lname, address bdate  
FROM project, department, employee  
WHERE dnum = pnumber AND mgrssn = ssn AND plocation = 'Stafford';
```

Qualifying Attribute with Relation Names and Aliasing

- Ambiguous attribute names:
 - **Attributes** in **different relations** can have the **same attribute name**.

- Example:

```
Works_on (Essn, Pno, Hours);  
Dependent (Essn, Dependent_name, Sex, Bdate, Relationship);
```

- When **Works_on** and **Dependent** relations appear in the **same query**, the attribute name **Essn** will be **ambiguous**
 - Find project numbers of projects worked on by employees who has a daughter names 'Alice'

```
SELECT Pno  
FROM Works_on, Dependent  
WHERE Essn = Essn AND name = 'Alice'
```

(You **cannot** tell that **Essn** is an attribute from the **Project** relation or from the **Dependent** relation

Solving the Ambiguous Attribute Names from Different Relations

- When different relations has a common attribute name, we can make that name un-ambiguous by:
 - **Qualifying** (= **prefixing**) the **attribute name** with the **source relation name**.

Project.essn	refers the the essn attribute in Project relation
Dependent.essn	refers the the essn attribute in Dependent relation

- Example:

- Find **project numbers** of **projects** worked on by **employees** who has a **daughter** named **'Alice'**.

```
SELECT pno
FROM   works_on, dependent
WHERE  works_on.essn = dependent.essn
AND    name = 'Alice';
```

Aliasing: qualifying attribute names from the same relation

- Fact:
 - Sometimes, we need to use the **same relation** **multiple times** in a **SELECT** command
 - Every **attribute name** of that **relation** will be **ambiguous**
- **Example:** for each employee, retrieve the employee's first and last name and the first and last name of his or her immediate supervisor.

```
SELECT  Employee.Fname, Employee.Lname, Employee.Fname,  
        Employee.Lname  
FROM    Employee, Employee  
WHERE   Employee.Super_ssn = Employee.Ssn;
```

Aliasing: qualifying attribute names from the same relation (cont.)

- **Alias**: solving ambiguous name from the **same relation**
- For each **employee**, retrieve the employee's first and last name and the first and last name of his or her immediate **supervisor**

```
SELECT  E.fname, E.lname, S.fname, S.lname  
FROM    Employee AS E, Employee AS S  
WHERE   E.superssn = S.ssn;
```

This practice is recommended since it results in queries that are easier to comprehend.

Renaming the Output Attribute Name

- Consider the **output**

```
SELECT e.fname, e.lname,
FROM employee e, emplo
WHERE e.superssn = m.ss
```

Output:

fname	lname
John	Smith
Ramesh	Narayan
Joyce	English
Alicia	Zelaya
Ahmad	Jabbar
Frankl	Wong
Jennif	Wallace

```
SELECT e.fname EmpFN, e.lname EmpLN, m.fname SupFN, m.lname SupFN
FROM employee e, employee m
WHERE e.superssn = m.ssn;
```

Output:

EmpFN	EmpLN	SupFN	SupFN
John	Smith	Frankl	Wong
Ramesh	Narayan	Frankl	Wong
Joyce	English	Frankl	Wong
Alicia	Zelaya	Jennif	Wallace
Ahmad	Jabbar	Jennif	Wallace
Frankl	Wong	James	Borg
Jennif	Wallace	James	Borg

The **name (title)** of the **attributes** are not very **meaningful**

The * Selector

- The **SELECT *** selects **all attributes** in all relations in the **FROM clause**

Q1C: **SELECT** *

FROM EMPLOYEE

WHERE Dno = 5;

Q1D: **SELECT** *

FROM EMPLOYEE, DEPARTMENT

WHERE Dname = 'Research' **AND** Dno = Dnumber;

Q10A: **SELECT** *

FROM EMPLOYEE, DEPARTMENT;

Tables as Sets in SQL

- SQL usually treats a table **not** as a set:
 - **Duplicate tuples** can **appear more than once** in a table, and in the result of a query
 - But an SQL table with a key is restricted to being a set
 - Use the **DISTINCT** keyword in the SELECT clause to **remove duplicate values**
 - Retrieve the salary of **every** employee and all **distinct** salary values

Q11: SELECT ALL Salary
 FROM EMPLOYEE;

Q11A: SELECT DISTINCT Salary
 FROM EMPLOYEE;

(a)	Salary	(b)	Salary
	30000		30000
	40000		40000
	25000		25000
	43000		43000
	38000		38000
	25000		55000
	25000		
	55000		

Tables as Sets in SQL (cont.)

- R

```
(SELECT pname
FROM   project, works_on, employee
WHERE  pnumber = pno
      AND  essn = ssn
      AND  lname = 'Smith')
UNION
(SELECT pnumber
FROM   project, works_on, employee
WHERE  pnumber = pno
      AND  essn = ssn
      AND  lname = 'Borg');
```

has the following set operators

+	+
	pname
+	+
	ProductX
	ProductY
	20
+	+

<----- How is 20 a pname ????

- Example: MySQL (version 5.5.24) will even accept the following UNION:
 - MySQL do not know about the meaning of each attribute

Other Tuple Conditions

- We have seen that you can use "ordinary" Boolean conditions in the where clause
- Other tuple conditions
 - The **membership** test operator (**IN** and **NOT IN**)
 - The **any selection** operator (**ANY**)
 - The **all selection** operator (**ALL**)
 - The **non-empty** set test operator (**EXISTS**)
 - The **null** test operator (**IS NULL**)
 - **Wild card** characters (**_** and **%**)

The IN and NOT-IN Comparison Operator

- The comparison operator **IN** tests **whether** a value **is contained** in a set
 - Syntax: attr **IN** (set of values)
- Example:
 - Find the fname and lname of employees whose SSN **is** 123456789 or 333445555
 - Find the fname and lname of employees whose SSN **is not equal to** 123456789 or 333445555

```
SELECT fname, lname
FROM   employee
WHERE  ssn IN ( ' 123456789', ' 333445555' );
```

```
SELECT fname, lname
FROM   employee
WHERE  ssn NOT IN ( ' 123456789', ' 333445555' );
```

The ANY Selection Operator

- ANT

- Syntax: attr RelationalOperator **any** (set of values)
- Meaning:

- The **(Boolean) expression** evaluates to **true** if the expression attr RelationalOperator x is **true** for **at least one member** of the (set of values)
- The **(Boolean) expression** evaluates to **false otherwise**

- Example:

```
SELECT fname, lname  
FROM employee  
WHERE salary >= ANY ( 30000, 50000 )
```

Note: why is the above query illegal?

- In the SQL syntax, you **cannot** use an **explicit set of value**
- Instead, you **must** use a **subquery** that **returns a set of values**

The ALL Selection Operator

- All

- Syntax: attr RelationalOperator **all** (set of values)

- Meaning:

- The **(Boolean) expression** evaluates to **true** if the expression attr RelationalOperator x is **true** for **all members** of the (set of values)
 - The **(Boolean) expression** evaluates to **false otherwise**

- Example

SELECT fname, lname

FROM employee

WHERE salary >= ALL (30000, 50000)

The != ALL Selection Operator

- Example: != **ALL**

```
SELECT fname, lname
```

```
FROM employee
```

```
WHERE ssn != ALL ('111-11-1111', '222-22-2222')
```

- x != **ALL** (set of values) is the same as x **NOT IN** (set of values)
- x = **ANY** (set of values) is the same as x **IN** (set of values)

Relations of One Tuple

- Example:
 - Find employees who salary is greater than or equal to John Smith

```
SELECT fname, lname, salary, dno
from employee A
where salary >= (select salary
                  from employee
                  where fname='John' and lname='Smith')
```

Because there is only **one** salary value for 'John Smith', we can **leave of ALL or ANY** in the set clause.

- Another one

```
SELECT fname, lname, salary, dno
from employee A
where salary >= (select salary
                  from employee
                  where dno = 5)
```

If you **omit** ALL or ANY keywords on a set of size ≥ 2 , SQL will report an **error** !!!

Because this query returns **more** than 1 tuple !

You get this error: Subquery returns more than 1 row

The EXISTSs and IS NULL

- Exists (set of values)

- Example:

- exists ('a', 'b') // returns true

- exists () // returns false

- IS NULL: testing for **null values**

- Find all employee tuples that has a NULL value in the salary attribute:

```
SELECT *  
FROM   employee  
WHERE  salary IS NULL
```

LIKE: Wildcard String Comparison

- Example: Find fname and lname of employees whose last name start with an 'S'?
- When **comparing strings**, you can use the **LIKE** operator to perform **wildcard string comparison**
 - Example: Find fname and lname of employees whose last name start with an 'S'
 - Solution:

```
SELECT fname, lname  
FROM   E.employee  
WHERE  lname LIKE 'S%'
```

Percent (%) matches 0 or more characters

LIKE: Wildcard String Comparison (cont.)

- Example: Find all employees who were born during the 1970s

- Solution:

```
SELECT Fname, Lname  
FROM EMPLOYEE  
WHERE Bdate LIKE '__ 7 _ _ _ _ _ _';
```

Underscore (_) matches exactly one character

- **Wildcards** can **ONLY** be used with the **LIKE** operator
 - The following query will find **people with last name equal to 'S%'**:

```
SELECT fname, lname  
FROM employee  
WHERE lname = 'S%'
```

More Tuple Conditions

- Another feature allows the use of **arithmetic** in queries
 - +, −, *, and division /
- **Example:** Show the resulting salaries if every employee working on the 'ProductX' project is given a 10% raise.

```
SELECT E.Fname, E.Lname, 1.1 * E.Salary AS Increased_sal  
FROM EMPLOYEE AS E, WORKS_ON AS W, PROJECT AS P  
WHERE E.Ssn = W.Essn AND W.Pno = P.Pnumber AND P.Pname = 'ProductX';
```

Nested Queries

- Nested query
 - A **nested query** is a **SELECT query** where the **WHERE clause** contains **subquery**.

```
SELECT ...  
FROM ...  
WHERE ... (SELECT ... /* Subquery */  
           FROM ...  
           WHERE ... )
```

- In theory, nesting can be arbitrarily deep

Examples of simple (= non-correlated) nested queries

- Query 1:
 - Find fname, lname of **employees** in the "Research" **department**
 - Solution: using a NON-nested query (uses a **join**):

```
SELECT fname, lname  
FROM   employee, department  
WHERE  dno = dnumber AND  dname = 'Research'
```

- Solution using a **nested query**

Examples of simple (= non-correlated) nested queries (cont.)

- Suppose we know that the 'Research' department has the department number 5, we can write the query as:

```
SELECT fname, lname  
FROM employee  
WHERE dno = 5
```

- We can use a **subquery** to find the **department number** of the **Research department**:

- (SELECT dnumber FROM department WHERE dname = 'Research')

- Therefore, the query as follows:

```
SELECT fname, lname  
FROM employee  
WHERE dno IN  
      ( SELECT pnumber  
        FROM department  
        WHERE dname = 'Research' )
```

Examples of simple (= non-correlated) nested queries (cont.)

- Query 2:

- Find fname, lname of **employees** that do not have any dependent

- 1) Conceptual solution

```
SELECT fname, lname
FROM   employee
WHERE  ssn IN
      { ssn of employees without dependent }
```

- 2) Which is equivalent to:

```
SELECT fname, lname
FROM   employee
WHERE  ssn NOT IN
      { ssn of employees with (one or more) dependent }
```

- 3) We can write this in SQL as:

```
SELECT fname, lname
FROM   employee
WHERE  ssn NOT IN
      (SELECT essn
       FROM   dependent);
```


Correlated Nested Queries


- So far, the nested queries above are **uncorrelated**:
 - The **inner (sub)query** is a **stand-alone** query that can be **executed independently** from the **outer query**
 - Example:

```
SELECT fname, lname  
FROM   employee  
WHERE  ssn NOT IN  
        (SELECT essn      /* Can be executed by itself */  
         FROM   dependent)
```

- The inner query is **completely independent** from the outer query.

Correlated Nested Queries (cont.)

- Correlated nested query
 - A query where the inner query (i.e., the query in the WHERE clause) **uses one or more attributes** from relation(s) specified in the outer query.



- The attribute attr4 of the relation R1 in the outer query is used in the inner query

Meaning of Ne

- Consider the following ex

```
SELECT fname, lname, salary, dno
FROM employee A
WHERE salary >= ALL (select salary
                    from emp
                    where B.dno=A.dno)
```

- Suppose the employee re
following tuples:

fname	lname	salary	dno
John	Smith	30000.00	5
Frankl	Wong	40000.00	5
Alicia	Zelaya	25000.00	4
Jennif	Wallace	43000.00	4
Ramesh	Narayan	38000.00	5
Joyce	English	25000.00	5
Ahmad	Jabbar	25000.00	4
James	Borg	55000.00	1

Outer tuple: A = John Smith 30000.00 5

```
WHERE salary >= ALL (select salary from employee B where B.dno=A.dno)
==> WHERE 30000.00 >= ALL (select salary from employee B where B.dno=5)
==> WHERE 30000.00 >= ALL (30000, 40000, 38000, 25000)
==> FALSE
```

Outer tuple: A = Frankl Wong 40000.00 5

```
WHERE salary >= ALL (select salary from employee B where B.dno=A.dno)
==> WHERE 40000.00 >= ALL (select salary from employee B where B.dno=5)
==> WHERE 40000.00 >= ALL (30000, 40000, 38000, 25000)
==> TRUE (select outer tuple !!!)
```

Outer tuple: A = Alicia Zelaya 25000.00 4

```
WHERE salary >= ALL (select salary from employee B where B.dno=A.dno)
==> WHERE 25000.00 >= ALL (select salary from employee B where B.dno=4)
==> WHERE 25000.00 >= ALL (25000, 43000)
==> FALSE
```

Output:

fname	lname	salary	dno
Frankl	Wong	40000.00	5
Jennif	Wallace	43000.00	4
James	Borg	55000.00	1

Examples on How to Use Correlated Nested Queries

- Example 1: Find fname, lname of **employees** that **do not have any dependent**

```
SELECT fname, lname
FROM Employee
WHERE { set of dependents of this employee }
      = empty set
```



```
SELECT fname, lname
FROM employee
WHERE ( select *
        from dependent
        where essn = employee.ssn)
      = empty set
```



```
SELECT fname, lname
FROM employee
WHERE not exists (select *
                   from dependent
                   where essn = employee.ssn)
```

Examples on how to use correlated nested queries (cont.)

- Example 2: Find fname, lname of **employees** that **do not** work on any **project** controlled by the **Research department**

```
SELECT fname, lname
FROM Employee A
WHERE { set of projects w.o. by employee A
      and controlled by Research dept }
      = empty set
```



```
SELECT fname, lname
FROM Employee A
WHERE not exists
      { set of projects w.o. by employee A
      and controlled by Research dept }
```



```
SELECT fname, lname
FROM employee A
WHERE not exists
      ( select *
        from projects
        where project w.o. by employee A
          and project controlled by Research dept )
```



```
SELECT fname, lname
FROM employee A
WHERE not exists
      (select *
       from projects
       where pnumber in ( project w.o. by employee A )
       and pnumber in ( project controlled by Research dept )
      )
```

Examples on how to use correlated nested queries (cont.)

- Example 2 (cont.): Find fname, lname of **employees** that **do not** work on any **project** controlled by the **Research department**

```
SELECT fname, lname
FROM employee A
WHERE not exists
    (select *
     from project
     where pnumber in (select pno
                       from works_on
                       where essn = A.ssn )

    and
        pnumber in (select pnumber
                    from project, department
                    where dnum = dnumber and  dname = 'Research' ) )
```

Ordering of Query Results

- ORDER BY clause
 - Order the tuples in the result of a query by the values of one or more of the attributes that appear in the query result
- Example:
 - Retrieve a list of **employees** and the **projects** they are **working on**, **ordered by department** and, **within each department, ordered alphabetically by last name, then first name**.

```
SELECT D.Dname, E.Lname, E.Fname, P.Pname
FROM DEPARTMENT AS D, EMPLOYEE AS E, WORKS_ON AS W, PROJECT AS P
WHERE D.Dnumber = E.Dno AND E.Ssn = W.Essn AND W.Pno = P.Pnumber
ORDER BY D.Dname, E.Lname, E.Fname;
```

INSERT, DELETE, and UPDATE Statements in SQL

- The insert command
 - The INSERT command can be used to **insert** one or more **tuples** into a **relation**
- Format 1: inserting one **complete** tuple:
 - Must specify **all** attribute values **in the exact order** at the **relation schema**

```
INSERT INTO    EMPLOYEE
VALUES        ( 'Richard', 'K', 'Marini', '653298653', '1962-12-30', '98
               Oak Forest, Katy, TX', 'M', 37000, '653298653', 4 );
```

- Format 2: inserting a **partial** tuple
 - Specify a **subset** of the attribute values in the **any order**

```
INSERT INTO employee(fname, lname, ssn)
VALUES ('Richard', 'Marini', '222669999');
```

NOTE: The attributes **not in the list** are given their **DEFAULT** or **NULL** value

INSERT Command in SQL

- DBMS should enforce all the integrity constraints due to insert command
 - How about the following insert operation

```
INSERT INTO      EMPLOYEE (Fname, Lname, Ssn, Dno)
VALUES           ('Robert', 'Hatcher', '980760540', 2);
```

(U2 is rejected if referential integrity checking is provided by DBMS.)

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

```
INSERT INTO      EMPLOYEE (Fname, Lname, Dno)
VALUES           ('Robert', 'Hatcher', 5);
```

(U2A is rejected if NOT NULL checking is provided by DBMS.)

INSERT Command in SQL

- Inserting a tuple using the result from a SELECT command
 - Syntax: INSERT INTO relationName (SELECT ...)
- Example:
 - To create a temporary table that has the employee last name, project name, and hours per week for each employee working on a project

```
CREATE TABLE WORKS_ON_INFO  
  ( Emp_name VARCHAR(15), Proj_name VARCHAR(15), Hours_per_week DECIMAL(3,1) );  
  
INSERT INTO WORKS_ON_INFO ( Emp_name, Proj_name, Hours_per_week )  
  SELECT  E.Lname, P.Pname, W.Hours  
  FROM    PROJECT P, WORKS_ON W, EMPLOYEE E  
  WHERE   P.Pnumber = W.Pno AND W.Essn = E.Ssn;
```

DELETE Command in SQL

- Delete one or more tuples from a relation using the DELETE command.
 - Syntax:

```
DELETE FROM relationName  
WHERE tuple-boolean-condition
```

- Example 1: Delete all **employees** from the 'Research' department

```
DELETE FROM Employee  
WHERE dno IN (SELECT dnumber  
              FROM department  
              WHERE dname = 'Research')
```

DELETE Command in SQL (cont.)

- Example 3:
 - Delete all employees from the 'Research' department have more than 2 dependents:

```
DELETE FROM employee
WHERE dno IN ( SELECT dnumber
                FROM department
                WHERE dname = 'Research')
AND ssn IN ( SELECT essn
              FROM dependent
              GROUP BY essn
              HAVING COUNT(name) > 2)
```

UPDATE Command in SQL

- Update one attribute value in one or more tuples from a relation using the UPDATE command

- Syntax:

```
UPDATE relationName  
SET  attributeName = expression  
WHERE tuple-boolean-condition
```

- Example1: Change the address of employee 'John Smith' to '123 Pike Lane, Austin, TX'

```
UPDATE employee  
SET  address = '123 Pike Lane, Austin, TX'  
WHERE fname = 'John' AND lname = 'Smith';
```

UPDATE Command in SQL (cont.)

- Example2: Give all employees in the 'Research' department a 10% raise

```
UPDATE Employee
SET      salary = 1.1 * salary
WHERE    dno IN ( SELECT dnumber
                  FROM  department
                  WHERE  dname = 'Research')
```