

# Chapter 7: More SQL

# Content

- More complex features of SQL retrieval queries
  - Nested queries aggregate functions, grouping
- Defining VIEWS on the database
- DROP and ALTER TABLE statement
- CREATE ASSERTION statement, CREATW TRIGGER statement

# Aggregate Functions in SQL

- Aggregate functions
  - **Compute** an **aggregate** value from a **set of values**
  - The **aggregate functions** available in SQL are:

1. **SUM()**: sums a set of values
2. **AVG()**: takes the average of a set of values
3. **MAX()**: finds the maximum value of a set of values
4. **MIN()**: finds the minimum value of a set of values
5. **COUNT()**: returns the number of elements in a set (**duplicates** are counted multiple times !!!)
6. **COUNT(DISTINCT ...)**: returns the number of **distinct** elements in a set (*duplicates are counted ONCE*)

# Applying Set of Function on a Set of Values

- Syntax:
  - SELECT AggFunc( attribute )
  - Apply the aggregate function **AggFunction** on the **selected set of tuples** returned by the **SELECT command**

```
SELECT salary  
FROM employee
```

```
SALARY  
30000  
40000  
25000  
43000  
38000  
25000  
25000  
55000
```

```
SELECT SUM(salary), AVG(salary), MAX(salary),  
       MIN(salary), COUNT(salary), COUNT( DISTINCT salary )  
FROM   employee
```

SUM(SALARY)	AVG(SALARY)	MAX(SALARY)	MIN(SALARY)	COUNT(SALARY)	COUNT(DISTINCT salary)
281000	35125	55000	25000	8	6

# Using the Result of Aggregate Functions in Queries

- Query
  - Find the fname and lname of the employee in the 'Research' department that earns **more than the average salary within the 'Research' department**

```
SELECT  fname, lname
FROM    employee
WHERE   salary > ( SELECT  AVG(salary)
                  FROM    employee, department
                  WHERE   dno = dnumber AND  dname = 'Research')
```

Correct?      No!

- Reason

- The  
de

```
SELECT  fname, lname
FROM    employee, department
WHERE   dno = dnumber      /* The employee must work for */
      AND  dname = 'Research' /* the Research department !! */
      AND  salary > (SELECT  AVG(salary)
                  FROM    employee, department
                  WHERE   dno = dnumber
                  AND     dname = 'Research')
```

of the Research

# Forming Groups with Common Attribute Values:

- **Grouping**

- Forming

- Each group has some attribute value in common
- Example

**Employee**

SSN	FName	other attributes	Sex	DNO	Salary
111-22-3333	John	.....	M	4	40000
123-45-6789	Mary	.....	F	5	50000
987-82-9823	James	.....	M	5	60000
982-71-9927	Jake	.....	M	4	50000

Group by DNO

111-22-3333	John	.....	M	4	40000
982-71-9927	Jake	.....	M	4	50000

123-45-6789	Mary	.....	F	5	50000
987-82-9823	James	.....	M	5	60000

partitionization


the value of

# Another Example:

## Employee

SSN	FName	other attributes	Sex	DNO	Salary
111-22-3333	John	•••••	M	4	40000
123-45-6789	Mary	•••••	F	5	50000
987-82-9823	James	•••••	M	5	60000
982-71-9927	Jake	•••••	M	4	50000

Group by (DNO, Sex)



Grouping attribute: dno and sex

# Forming Groups with Common Attribute Values: Grouping Attributes (cont.)

- **GROUP BY** clause for this purpose

Here is the employee relation:

```
SELECT fname, lname, sex, dno, salary
FROM employee
```

fname	lname	sex	dno	salary
James	Borg	M	1	55000.00
Alicia	Zelaya	F	4	25000.00
Jennif	Wallace	F	4	43000.00
Ahmad	Jabbar	M	4	25000.00
John	Smith	M	5	30000.00
Frankl	Wong	M	5	40000.00
Ramesh	Narayan	M	5	38000.00
Joyce	English	F	5	25000.00

- Example

- Find the total salary paid in each department

Here is the solution:

```
SELECT dno, sum(salary)
FROM employee
GROUP BY dno
```

DNO	SUM(SALARY)
1	55000
4	93000
5	133000



# Grouping with Multiple Common Attribute Values

- Query:
  - Find the total salary paid to **male** and **female** employees (separate total) for **each department**
- Solution:

```
SELECT    dno, sex, sum(salary)
FROM      employee
GROUP BY  dno, sex
```

DNO	S	SUM(SALARY)
1	M	55000
4	F	68000
4	M	25000
5	F	25000
5	M	108000

# Note on Output Format

- See an example first:

```
SUM(SALARY)
```

```
-----
```

```
25000
```

```
55000
```

```
108000
```

```
68000
```

```
25000
```

```
SELECT      sum(salary)  /* Omit dno, sex */  
FROM        employee  
GROUP BY    dno, sex
```

What does each row mean?  
Group by What ?

You don't have a clue what  
each row mean...

- You don't **need** to specify the **grouping attributes** in the **SELECT clause**
- **However**, if you **omit** , you will have **no idea what** the **result mean** !!!

# Execution of an SQL Query with GROUP BY Clause

- A query with a GROUP BY clause is processed as follows:
  - First, **select** the **tuples** that satisfies the **tuple (WHERE) condition**
  - Then, the **selected tuples** in step 1 are **grouped** based on their value in the **grouping attributes**
  - Finally, one or more aggregate functions is applied to the groups
- Example:

How is the following query processed?

```
SELECT    dno, sum(salary)
FROM      employee
WHERE     sex = 'M'
GROUP BY  dno
```

The complete Employee relation:

SSN	DNO	S	SALARY
123456789	5	M	30000
333445555	5	M	40000
999887777	4	F	25000
987654321	4	M	43000
666884444	5	M	38000
453453453	5	F	25000
987987987	4	M	25000
888665555	1	M	55000

(1) The **WHERE** clause is processed **first**:

SSN	DNO	S	SALARY
123456789	5	M	30000
333445555	5	M	40000
987654321	4	M	43000
666884444	5	M	38000
987987987	4	M	25000
888665555	1	M	
55000aaaa			

Output:

DNO	SUM(SALARY)
1	55000
4	68000
5	108000

(2) Resulting tuples are **grouped** by the **grouping attributes (dno)**:

SSN	DNO	S	SALARY
888665555	1	M	55000
987654321	4	M	43000
987987987	4	M	25000
123456789	5	M	30000
333445555	5	M	40000
666884444	5	M	38000

(3) Apply the **aggregate function(s)** on each group

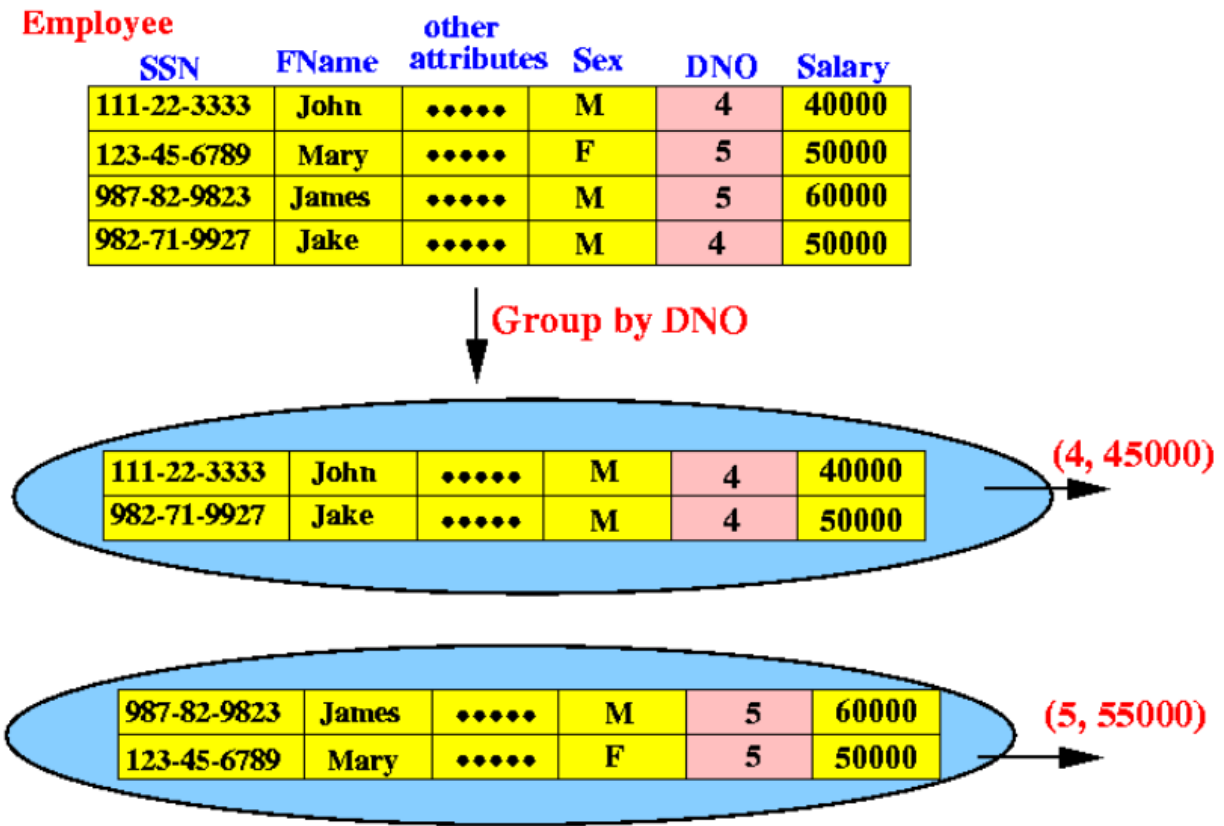
-----> (DNO=1) SUM = 55000

-----> (DNO=4) SUM = 68000

-----> (DNO=5) SUM = 108000

# Effect of the GROUP BY on Common Attribute Values

- Property:
  - After forming groups based on the **grouping attributes**, the **attribute values other than the grouping attributes** will be **lost**



# Effect of the GROUP BY on Common Attribute Values (cont.)

- Why?

- The **group by** function will **produce one (1) tuple** output for *each group*

Therefore, you can *only* have **a single value** as **output**

- *Each tuple* in a **group** has the **same value** for **DNO** (**grouping attribute**)

Therefore, you *can* produce **a single value** for **DNO** as an **output value**

- *Different tuples* in the **same group** have **different value** for the **other attributes** (e.g., **SSN** in **group (DNO=4)**: 111-22-3333 and 982-71-9927)

Therefore, you *cannot* produce **a single value** for **non-grouping attributes** as an **output value**  
!!!

# Examples of GROUP BY Queries

- Query 1:

- For
- and

- Solution

```
SELECT
FROM
WHERE
```

**CORRECT:**

```
SELECT  dname, COUNT(ssn), AVG(salary)
FROM    employee, department
WHERE   dno = dnumber
GROUP BY dno, dname
```

of employees

salary)

ssn			
12345678			
33344555			
99988777			
98765432			
66688444			
45345345			
987987987	4	Administration	25000.00
888665555	1	Headquarters	55000.00

DNAME	COUNT (SSN)	AVG (SALARY)
Administration	3	31000
Headquarters	1	55000
Research	4	33250

Wrong! dname must be a grouping attribute !!

# Flashback: Data Independence

- Physical data independence

▪ **Physical data independence** = the **ability** to **change** the **Physical storage format** of the **data files** *without* having to **change** the **program** (so you don't need to *re-compile* the code !!!)

- You can **change** the **structure of the physical data records** **without** having to **change the programs** to **access the new record structure**
- All you need to do is: **provide** a **new (updated) data description** (= meta data)

- Logical data independence

▪ **Logical data independence** = the **ability** to **present** the **stored data** in **any format** to the **user**

- Data is stored in **one format** (called the **conceptual database schema**), the database system can **present the data in different ways** to different users
- The view of the data can be adapted to the need of the user



# Concept of a View in SQL

- A **view** is a **virtual relation** that is **derived** from
  - Relations
  - Other **virtual relations**
- Property of a virtual relation:
  - A virtual relation (= view) does **not exist** as in **physical form**
  - The tuples in a virtual relation (= view) **are not stored** physically in the database
  - The tuples in a virtual relation (= view) are computed as a **temporary relation** when the virtual relation is used

- the **virtual relation** is **constructed** as a **temporary relation** at **that moment** in time

# Example Temporary Rel

- Suppose a secretary maintains a **list**
  - The list contains the following inform

FName	LName	ProjectName	#HoursWork
-------	-------	-------------	------------

- You could **obtain the data** using this S

```
SELECT fname, lname, pname, hours
FROM   employee, works_on, project
WHERE  ssn = essn
      AND pno = pnumber
```

Output:

FNAME	LNAME	PNAME	HOURS
John	Smith	ProductX	32.5
John	Smith	ProductY	7.5
Frank	Wong	ProductY	10
Frank	Wong	Reorganization	10
Frank	Wong	ProductZ	10
Frank	Wong	Computerization	10
Joyce	English	ProductX	20
Joyce	English	ProductY	20
John	Doe	ProductZ	40
James	Borg	Reorganization	0
Jack	Wallace	Reorganization	15
Jack	Wallace	Newbenefits	20
Jake	Jones	Computerization	35
Jake	Jones	Newbenefits	5
Alice	Miller	Computerization	10
Alice	Miller	Newbenefits	30

Question: How to find the activities of "John Smith"?

# Example Temporary Relation (cont.)

- To find the **activities** of "John Smith", we can use the following **temporary relation** construct:

```
SELECT *  
FROM ( SELECT fname, lname, pname, hours  
      FROM employee, works_on, project  
      WHERE ssn = essn  
            AND pno = pnumber  
      ) EmpActivity  
WHERE fname = 'John'  
      AND lname = 'Smith'
```

Output:

FNAME	LNAME	PNAME	HOURS
John	Smith	ProductX	32.5
John	Smith	ProductY	7.5

The **temporal relation**  
EmpActivity is created on  
the fly !!!!

# Creating View (Virtual Relations)

- A **view** (= **virtual relation**) is defined using the **CREATE VIEW** command
  - A (virtual) table name (or **view name**)
  - **A list of attribute names**, and
  - A **query** to **specify the contents** of the view.
- Example:

```
CREATE VIEW EmpActivity
AS
( SELECT Fname, Lname, Pname, Hours
  FROM employee, works_on, project
 WHERE ssn = essn
       AND pno = pnumber
 )
```

EmpActivity

Fname	Lname	Pname	Hours
-------	-------	-------	-------

Results:

- The **CREATE VIEW** command defines an **virtual table (relation)** called EmpActivity
- This **virtual table/relation** will be **construct** when we use the EmpActivity relation in a **query !!!**

# Creating View (Virtual Relations) (Example)

- We can now **specify SQL queries** on a **view**—or virtual **table**—in the **same way** we specify queries involving **base tables**
  - Example: Find all activities of John Smith

```
SELECT *  
FROM   EmpActivity  
WHERE  fname = 'John'  
       AND  lname = 'Smith'
```

FNAME	LNAME	PNAME	HOURS
John	Smith	ProductX	32.5
John	Smith	ProductY	7.5

# Why Do We Introduce t

- Now thinking: **why** do we introduce the concept of **view**?
  - **Simplify** the **specification** of certain queries
  - Used as a **security and authorization mechanism**

```
SELECT fname, lname, pname, hours
FROM   employee, works_on, project
WHERE  ssn = essn
      AND pno = pnumber
```

Output:

FNAME	LNAME	PNAME	HOURS
John	Smith	ProductX	32.5
John	Smith	ProductY	7.5
Frank	Wong	ProductY	10
Frank	Wong	Reorganization	10
Frank	Wong	ProductZ	10
Frank	Wong	Computerization	10
Joyce	English	ProductX	20
Joyce	English	ProductY	20
John	Doe	ProductZ	40
James	Borg	Reorganization	0
Jack	Wallace	Reorganization	15
Jack	Wallace	Newbenefits	20
Jake	Jones	Computerization	35
Jake	Jones	Newbenefits	5
Alice	Miller	Computerization	10
Alice	Miller	Newbenefits	30

# Views as Authorization Mechanism

- **Views** can be used to **hide** certain **attributes or tuples** from **unauthorized users**
  - Example: suppose a certain user is **only allowed** to see **employee** information for employees who **work for department 5**;

```
CREATE VIEW DEPT5EMP
AS
  ( SELECT *
    FROM EMPLOYEE
    WHERE Dno = 5;
  )
```

Grant the user the privilege to query the **view** **but not** the **base table EMPLOYEE** itself.

- Another example: restrict a user to only see certain columns: first name, last name, and address of an employee

```
CREATE VIEW BASIC_EMP_DATA
AS
  ( SELECT Fname, Lname, Address
    FROM EMPLOYEE;
  )
```

# Keep View be Up-to-Date

- A view **is supposed to** be always **up-to-date**
- If we **modify** the tuples in the **base tables** on which the **view is defined**
  - The view must automatically **reflect these changes**
  - The view **does not have to be** realized or materialized **at the time of view definition but rather at the time when we specify a query on the view**
  - Responsibility of the DBMS and **not** the user

```
CREATE VIEW EmpActivity
AS
( SELECT fname, lname, pname, hours
  FROM employee, works_on, project
 WHERE ssn = essn
       AND pno = pnumber
 )
```



# View is Computed on the fly

- Consider the following **view definition**

```
CREATE VIEW Dept_Info(dname, no_emps, total_sal)
AS
  ( SELECT      dname, count(ssn), sum(salary)
    FROM      department, employee
   WHERE      dnumber = dno
   GROUP BY   dname
  )
```

```
SELECT *
FROM   Dept_Info
WHERE  no_emps > 2
```

DNAME	NO_EMPS	TOTAL_SAL
Administration	3	93000
Research	4	133000

- Example query on this view:
  - Find all department with >2 employees:

# View is Computed on the fly (cont.)

- Change the salary of employees:

```
UPDATE employee  
SET    salary = salary + 10000
```

```
SELECT *  
FROM   Dept_Info  
WHERE  no_ems > 2
```

DNAME	NO_EMPS	TOTAL_SAL
Research	4	173000
Administration	3	123000

# Modifying Existing Relations: Altering and Dropping

- Modifying the *structure* of an existing relation
  - **After** you have **defined a relation** using the CREATE TABLE command, you can still **make changes** to its **structure**
- **DROP TABLE**
  - **Deleting the definition of a table** from the **database schema**
  - Syntax: DROP TABLE relation\_name
  - The **relation definition** is **removed**, along with **any tuples (data)** that you have previously inserted....
  - Here is the **difference** between **Drop** and **Delete**

# Modifying Existing Relations: Altering and Dropping (cont.)

- Two types:
  - Example: if we no longer wish to keep track of dependents of employees in the COMPANY database, drop it:
  - **1) DROP TABLE DEPENDENT CASCADE;**
    - All **constraints, views, and other elements** that **reference the table** being dropped are **also dropped automatically** from the schema, along with the table itself.
  - **2) DROP TABLE DEPENDENT RESTRICT;**
    - A table is dropped only if it is **not referenced in any constraints** (for example, by foreign key definitions in another relation) or **views** (see Section 7.3) or by **any other elements**

# Modifying Existing Relations: Altering and Dropping (cont.)

- **ALTER TABLE**: changing a relation schema.
  - SQL allows the owner of the database relation to change it by:

- |   |
|---|
| <ul style="list-style-type: none"><li>▪ <b>adding</b> one or more <b>attributes</b> to the relation</li><li>▪ <b>removing</b> one or more <b>attributes</b> from the relation</li></ul>   |
| <ul style="list-style-type: none"><li>▪ <b>adding</b> one or more <b>constraints</b> to the relation</li><li>▪ <b>removing</b> one or more <b>constraints</b> from the relation</li></ul> |

- **ALTER TABLE** COMPANY.EMPLOYEE **ADD COLUMN** Job VARCHAR(12);
- **ALTER TABLE** COMPANY.EMPLOYEE **DROP COLUMN** Address CASCADE;

# Create Assertion

- **CREATE ASSERTION**

- Used to **specify additional types of constraints** that are **outside the scope** of the **built-in** relational model **constraints**
- Example: **specify the constraint:** *the salary of an employee **must not be greater than** the salary of the manager of the department that the employee works for:*

```
CREATE ASSERTION SALARY_CONSTRAINT
CHECK ( NOT EXISTS ( SELECT      *
                      FROM        EMPLOYEE E, EMPLOYEE M,
                                DEPARTMENT D
                      WHERE        E.Salary>M.Salary
                                AND   E.Dno = D.Dnumber
                                AND   D.Mgr_ssn = M.Ssn ) );
```

# Create Trigger

- **CREATE TRIGGER**

- Used to specify **automatic actions** that the database system will **perform** when certain **events and conditions occur**.

- Example:

- *Suppose we want to check whenever an employee's salary is greater than the salary of his or her direct supervisor*
  - **Trigger conditions:** inserting a new employee record, changing an employee's salary, or changing an employee's supervisor
  - **Action:** suppose that the action to take would be to call an external stored procedure SALARY\_VIOLATION, which will notify the supervisor.

# Create Trigger (cont.)

- The trigger could then be written as

```
CREATE TRIGGER                                SALARY_VIOLATION  
BEFORE INSERT OR UPDATE OF  SALARY, SUPERVISOR_SSN ON EMPLOYEE  
FOR EACH ROW  
    WHEN ( NEW.SALARY > ( SELECT SALARY FROM EMPLOYEE  
                           WHERE SSN = NEW.SUPERVISOR_SSN  
                           )  
    )  
    INFORM_SUPERVISOR(NEW.Supervisor_ssn, NEW.Ssn );
```

**ECA (Event, Condition, Action)**